

Class: Final Year (Computer Science and Engineering)

Year: 2021-22 **Semester:** 1

Course: High Performance Computing lab

ESE Exam

22/11/2021 01.00 PM - 04.00 PM

Exam Seat No:2018BTECS00100

Name:Prakash Singh

Exam Seat Number:2018BTECS00100

Problem Statement 1

Statement 1:Write an OpenMP program to print inverted pyramid using *.

Screenshot #:

```
prax@praxx-ideapad:~/Desktop/HPC_ESE$ gcc -fopenmp 1.c
prax@praxx-ideapad:~/Desktop/HPC_ESE$ ./a.out
* * * * *
 * * * * *
  * * * *
   * * *
    * *
     *

Executed Done in 0.000027 seconds
prax@praxx-ideapad:~/Desktop/HPC_ESE$
```

```
int main(int argc, char*argv[]) {
    int rows=N, i, j, space;

    double time = omp_get_wtime();
    #pragma omp for
    for (i = rows; i >= 1; --i) {
        for (space = 0; space < rows - i; ++space)
            printf(" ");
        for (j = i; j <= 2 * i - 1; ++j)
            printf("* ");
        for (j = 0; j < i - 1; ++j)
            printf("* ");
        printf("\n");
    }

    printf("\nExecuted Done in %f seconds\n", omp_get_wtime()-time);
    return 0;
}
```

Information #:Used openMP to print inverted pyramid using *

Problem Statement 2

Statement 2:Implement MPI program to reduce the data from n processes to root process.

Screenshot #:

```
prax@praxx-ideapad:~/Desktop/HPC_ESE$ mpicc 2 2.c
prax@praxx-ideapad:~/Desktop/HPC_ESE$ mpiexec -n 10 ./2
[Process 4]: has local data 9
[Process 5]: has local data 11
[Process 1]: has local data 3
[Process 2]: has local data 5
[Process 6]: has local data 13
[Process 0]: has local data 1
[Process 7]: has local data 15
[Process 3]: has local data 7
[root Process 0]: has the result: 100
[Process 8]: has local data 17
[Process 9]: has local data 19
prax@praxx-ideapad:~/Desktop/HPC_ESE$
```

```
#include <stdio.h>
#include "mpi.h"

int main(int argc, char **argv) {
    int size, rank;
    int localdata;
    int result;
    const int root = 0;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    //Prepare local data
    localdata = rank*2+1;
    printf("[Process %d]: has local data %d\n", rank, localdata);

    //MPI_Reduce(sendbuf, recvbuf, count, MPI_INT, MPI_SUM, root, MPI_COMM_WORLD);
    MPI_Reduce(&localdata, &result, 1, MPI_INT, MPI_SUM, root, MPI_COMM_WORLD);

    if (rank == root) {
        printf("[root Process %d]: has the result: %d", rank, result);
        printf("\n");
    }

    MPI_Finalize();
    return 0;
}
```

Information #:Used MPI_reduce function to reduce data from n process to the root process

Statement 3: Implement Matrix-Vector multiplication using CUDA.

Screenshot :

```
prax@praxx-ideapad:~/Desktop/HPC_ESE$ nvcc 3.cu
prax@praxx-ideapad:~/Desktop/HPC_ESE$ ./a.out

Vector:
2      2      2      2      2      2      2      2      2      2
Matrix:
0      0      0      0      0      0      0      0      0      0
1      1      1      1      1      1      1      1      1      1
2      2      2      2      2      2      2      2      2      2
3      3      3      3      3      3      3      3      3      3
4      4      4      4      4      4      4      4      4      4
5      5      5      5      5      5      5      5      5      5
6      6      6      6      6      6      6      6      6      6
7      7      7      7      7      7      7      7      7      7
8      8      8      8      8      8      8      8      8      8
9      9      9      9      9      9      9      9      9      9

Product of vector and matrix:
0
20
40
60
80
100
120
140
160
180
prax@praxx-ideapad:~/Desktop/HPC_ESE$
```

```
#include<stdio.h>
#include<cuda.h>
#define row1 10
#define col1 10
#define col2 10

__global__ void matrix_multiply(int *l,int *m, int *n)
{
    int x=blockIdx.x;
    int y=blockIdx.y;
    int k;

    n[col2*y+x]=0;
    for(k=0;k<col1;k++)
    {
        n[col2*y+x]=n[col2*y+x]+l[col1*y+k]*m[col2*k+x];
    }
}

int main()
{
    int a[row1][col1];
    int b[col2];
    int c[row1][col2];
    int *d,*e,*f;
    int i,j;

    for(i=0;i<row1;i++)
    {
        for(j=0;j<col1;j++)
        {
            a[i][j]=i;
        }
    }

    for(i=0;i<col2;i++)
```

```
        b[i]=2;
    }
    cudaMalloc((void **)&d,row1*col1*sizeof(int));
    cudaMalloc((void **)&e,col2*sizeof(int));
    cudaMalloc((void **)&f,row1*col2*sizeof(int));

    cudaMemcpy(d,a,row1*col1*sizeof(int),cudaMemcpyHostToDevice);
    cudaMemcpy(e,b,col2*sizeof(int),cudaMemcpyHostToDevice);

    dim3 grid(col2,row1);
    matrix_multiply<<<grid,1>>>(d,e,f);

    cudaMemcpy(c,f,row1*col2*sizeof(int),cudaMemcpyDeviceToHost);

    printf("\nVector:\n");
    for(i=0;i<col2;i++)
        printf("%d\t",b[i]);

    printf("\nMatrix:\n");
    for(i=0;i<row1;i++)
    {
        for(j=0;j<col1;j++)
        {
            printf("%d\t",a[i][j]);
        }
        printf("\n");
    }
    printf("\nProduct of vector and matrix:\n ");
    int sum=0;
    for(i=0;i<row1;i++)
    {
        sum=0;
        for(j=0;j<col2;j++)
        {
            sum+=c[i][j];
        }
        printf("%d\t",sum);
        printf("\n");
    }
    cudaFree(d);
    cudaFree(e);
    cudaFree(f);

    return 0;
```

Technologies Used: OpenMP, MPI, CUDA

GitHub Link:https://github.com/prakx1/HPC_ESE